

# Lab 5 – Containers and Kubernetes

## Basics:

- Introduction to containers, e.g. Docker: <https://docs.docker.com/get-started/>
  - Lightweight virtualization
  - Key concepts:
    - Dockerfile, Image, Container, Image repositories
- Kubernetes, container orchestration: <https://kubernetes.io/docs/concepts/overview/what-is-kubernetes/>
  - Cluster, Pods, Services etc.
  - Used to manage a cluster of hosts/containers
  - Why do we need orchestration? What functionalities are provided by k8s, which exceed those of a simple dockerd daemon?
- Alternative implementations:
  - Containers: apptainer, podman, lxc
  - Container orchestration: OpenShift, Rancher, Nomad, commercial services: e.g. AWS ECS, Azure AKS, Google GKE

## Prerequisites

- Preferably a Linux environment (Linux VM, or Windows/WSL 2)
- Docker
- Kubernetes cluster. While installing a full-fledged cluster is quite complicated, it's possible to prepare a development "mini k8s cluster" with the help of
  - Minikube: <https://minikube.sigs.k8s.io/docs/start> (simple, only 1-node cluster)
  - Kind: <https://kind.sigs.k8s.io> (Kubernetes in Docker, possible to emulate multiple nodes on a single machine)
  - Or others...

## Assignments

Assignments were performed in **WSL** in the **Ubuntu** operating system

### 1) Dockerize AWS-CLI (3p)

- a) Create a Dockerfile with [aws-cli](#), built from [source files](#) (using tar.gz file / make).
  - Use a [multi-stage build](#) to create a small image. Also use a [small base image, e.g., based on Alpine Linux](#).
  - Read more on [Python and Docker multistage build](#).
  - Follow [best practices for writing Dockerfiles](#). In particular read about [ADD or COPY](#).

nano Dockerfile

```
FROM python:3.8-alpine3.14 AS builder

ENV AWSCLI_VERSION=2.10.1

WORKDIR /aws-cli-docker

RUN apk add --no-cache \
    curl \
    make \
    cmake \
    gcc \
    g++ \
    libc-dev \
    libffi-dev \
    openssl-dev \
    && curl https://awscli.amazonaws.com/awscli-${AWSCLI_VERSION}.tar.gz | tar -xz \
    && cd awscli-${AWSCLI_VERSION} \
    && ./configure --prefix=/opt/aws-cli/ --with-download-deps \
    && make \
    && make install

FROM python:3.8-alpine3.14

COPY --from=builder /opt/aws-cli/ /opt/aws-cli

ENTRYPOINT ["/opt/aws-cli/bin/aws"]
```

My main example for creating the image was the [AWS instructions](#). To improve the Dockerfile, I concretized the Alpine version and removed the groff installation

- b) Build an image based on your Dockerfile and test it.
  - Use [a volume](#) (-v option) to make AWS credentials available in the container.

Running image locally

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ docker build -t aws-cli-docker .
[+] Building 10.2s (8/8) FINISHED
=> [internal] load build definition from Dockerfile
=> => transferring dockerfile: 596B
=> [internal] load metadata for docker.io/library/python:3.8-alpine3.14
=> [internal] load .dockerignore
=> => transferring context: 2B
=> [builder 1/3] FROM docker.io/library/python:3.8-alpine3.14@sha256:9688d36309a9d72d346ba5d1e9af94cd129fa9c6a985e6dfdd089536116b2db1
=> => resolve docker.io/library/python:3.8-alpine3.14@sha256:9688d36309a9d72d346ba5d1e9af94cd129fa9c6a985e6dfdd089536116b2db1
=> CACHED [builder 2/3] WORKDIR /aws-cli-docker
=> CACHED [builder 3/3] RUN apk add --no-cache curl make cmake gcc g++ libc-dev libffi-dev openssl-dev
=> CACHED [stage-1 2/2] COPY --from=builder /opt/aws-cli/ /opt/aws-cli
=> exporting to image
=> => exporting layers
=> => exporting manifest sha256:0f8dccad78e0ef01d40613c4d1f33b15057221c16b6d3996a9ebf08ab606c99c
=> => exporting config sha256:4c630e5a836c6abd6518de891b9e60ffc322252c4dc3000baaf339ad9ae30f2c
=> => exporting attestation manifest sha256:14f02fa43cb4ab09197ae9807a9401bdd233c9030c40d754ce88bc8ccb59729f
=> => exporting manifest list sha256:6fa473989f153b65b9a1b1affbf1ab7602345a83e0a893618d08ce275088666
=> => naming to docker.io/library/aws-cli-docker:latest
=> => unpacking to docker.io/library/aws-cli-docker:latest
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ docker run --rm -it aws-cli-docker --version
aws-cli/2.10.1 Python/3.8.13 Linux/5.15.153.1-microsoft-standard-WSL2 source-sandbox/x86_64.alpine.3 prompt/off
```

## Checking buckets on AWS with container

```
ubuntu@LAPTOP-9M7MBHQ5:~$ docker run --rm -it -v ~/.aws:/root/.aws aws-cli-docker --version
aws-cli/2.10.1 Python/3.8.13 Linux/5.15.153.1-microsoft-standard-WSL2 source-sandbox/x86_64.alpine.3 prompt/off
```

## Size of the container image

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ docker images
REPOSITORY          TAG         IMAGE ID      CREATED       SIZE
aws-cli-docker      latest      6fa473989f15  9 hours ago  487MB
```

## 2) Kubernetes deployment (5p)

### a) Create a k8s cluster using Amazon Elastic Kubernetes Service (EKS)

- You can also use minikube, kind, or any other Kubernetes distribution, or existing cluster.
- Minikube, by default, uses its own internal Docker daemon. This daemon doesn't know anything about images built previously. Prepare your environment by directing it to access the internal docker daemon by using the `$(minikube docker-env)` command and rebuild your images. This way images will be available within the k8s cluster.

(<https://medium.com/bb-tutorials-and-thoughts/how-to-use-own-local-doker-images-with-minikube-2c1ed0b0968>)

My main example was the [AWS instructions](#) for creating cluster with AWS console

## Cluster

The screenshot displays the AWS Management Console interface for an Amazon Elastic Kubernetes Service (EKS) cluster. The left sidebar shows the navigation menu with 'Amazon Elastic Kubernetes Service' selected. The main content area shows the 'Lab5' cluster details. The 'Cluster info' section indicates the cluster is in a 'Creating' status, using Kubernetes version 1.31, with standard support until November 26, 2025. The 'Details' section provides further information: the API server endpoint is '-', the OpenID Connect provider URL is '-', the certificate authority is '-', and the cluster IAM role ARN is 'arn:aws:iam::112106458942:role/LabRole'. The cluster was created 'a minute ago' and has a cluster ARN of 'arn:aws:eks:us-east-1:112106458942:cluster/Lab5'. The platform version is 'eks.6'.

Cluster info			
Status	Kubernetes version	Support period	Provider
Creating	1.31	Standard support until November 26, 2025	EKS

Details		
API server endpoint	OpenID Connect provider URL	Created
-	-	a minute ago
Certificate authority	Cluster IAM role ARN	Cluster ARN
-	arn:aws:iam::112106458942:role/LabRole	arn:aws:eks:us-east-1:112106458942:cluster/Lab5
	<a href="#">View in IAM</a>	Platform version
		eks.6

Nodes group

Group

↻

Edit











Delete

Node group configuration [Info](#)


Kubernetes version	AMI type <a href="#">Info</a>	Status
1.31	AL2_x86_64	✔ Active
AMI release version <a href="#">Info</a>	Instance types	Disk size
1.31.0-20241024	t3.medium	20 GiB

- Details
- Nodes
- Health issues 0
- Kubernetes labels
- Update config
- Kubernetes taints
- Update history
- Tags



Details

Node group ARN  <a href="#">arn:aws:eks:us-east-1:112106458942:nodegroup/Lab5/Group/18c9809a-000b-1cc8-5575-f6320245cbf2</a>	Autoscaling group name <a href="#">eks-Group-18c9809a-000b-1cc8-5575-f6320245cbf2</a> 	Capacity type On-Demand	Subnets <a href="#">subnet-09c8079647cd19442</a>  <a href="#">subnet-0cc365850abc66759</a>  <a href="#">subnet-059f9cf9d07bf534f</a>  <a href="#">subnet-0bdc5157bae1751c3</a>  <a href="#">subnet-09aa55da60921eb8d</a> 
Created  4 hours ago	Node IAM role ARN  <a href="#">arn:aws:iam::112106458942:role/LabRole</a> <a href="#">View in IAM</a> 	Desired size 2 nodes	Configure remote access to nodes off
		Minimum size 2 nodes	
		Maximum size 2 nodes	

Nodes (2) [Info](#)



< 1 >

Node name	Instance type	Node group	Created	Status
<a href="#">ip-172-31-13-24.ec2.internal</a>	t3.medium	<a href="#">Group</a>	Created  4 hours ago	✔ Ready
<a href="#">ip-172-31-76-228.ec2.internal</a>	t3.medium	<a href="#">Group</a>	Created  4 hours ago	✔ Ready

Node groups (1) [Info](#)

Edit

Delete

Add node group

	Group name	Desired size	AMI release version	Launch template	Status
<input type="radio"/>	<a href="#">Group</a>	2	1.31.0-20241024	-	✔ Active

## Cluster authentication

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ aws sts get-caller-identity
{
  "UserId": "AROARUGQQZ47DJK600WX0:user3610452=ayahorava@student.agh.edu.pl",
  "Account": "112106458942",
  "Arn": "arn:aws:sts::112106458942:assumed-role/voclabs/user3610452=ayahorava@student.agh.edu.pl"
}
```

## Create a kubeconfig file for the cluster

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ aws eks update-kubeconfig --region us-east-1 --name Lab5
Added new context arn:aws:eks:us-east-1:112106458942:cluster/Lab5 to /home/ubuntu/.kube/config
```

## Testing configuration

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ kubectl get svc
NAME                TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
kubernetes          ClusterIP     10.100.0.1    <none>         443/TCP          8m7s
```

- b) Using Helm, install an [NFS server and provisioner](#) in the cluster.
- Go to charts/nfs-server-provisioner for a README.
  - Pay attention to configuration parameters, in particular, override `storageClass.name` which denotes the name of the StorageClass that you'll have to use when creating Persistent Volume Claims.

## Installing Helm

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash
% Total    % Received % Xferd  Average Speed   Time    Time     Current
                                 Dload  Upload   Total   Spent    Left   Speed
100 11903  100 11903    0     0  113k      0  --:--:-- --:--:-- --:--:--  113k
Downloading https://get.helm.sh/helm-v3.16.2-linux-amd64.tar.gz
Verifying checksum... Done.
Preparing to install helm into /usr/local/bin
helm installed into /usr/local/bin/helm
```

## Installing nfs

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ helm repo add nfs-ganesha-server-and-external-provisioner https://kubernetes-sigs.github.io/nfs-ganesha-server-and-external-provisioner/
"nfs-ganesha-server-and-external-provisioner" has been added to your repositories
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ helm install my-release nfs-ganesha-server-and-external-provisioner/nfs-server-provisioner
NAME: my-release
LAST DEPLOYED: Tue Nov 5 20:18:50 2024
NAMESPACE: default
STATUS: deployed
REVISION: 1
TEST SUITE: None
NOTES:
The NFS Provisioner service has now been installed.

A storage class named 'nfs' has now been created
and is available to provision dynamic volumes.

You can use this storageclass by creating a `PersistentVolumeClaim` with the
correct storageClassName attribute. For example:
```

```
---
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: test-dynamic-volume-claim
spec:
  storageClassName: "nfs"
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 100Mi
```

- c) Create a [Persistent Volume Claim](#) which will bind to a NFS Persistent Volume [provisioned dynamically](#) by the provisioner installed in the previous step.

nano nfs-pvc.yaml

```
GNU nano 7.2 nfs-pvc.yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  storageClassName: "nfs"
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 100Mi
```

Running nfs-pvc.yaml

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ kubectl apply -f nfs-pvc.yaml
persistentvolumeclaim/my-pvc created
```

Checking pvc

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ kubectl get pvc
```

NAME	STATUS	VOLUME	CAPACITY	ACCESS MODES	STORAGECLASS	VOLUMEATTRIBUTESCLASS	AGE
my-pvc	Bound	pvc-92e684d9-8c83-4abb-9dab-9d31e6c74c2a	100Mi	RWO	nfs	<unset>	9s

- d) Create a [Deployment](#) with a HTTP server (e.g., apache or nginx). The web content directory should be mounted as a volume using the PVC created in the previous step.

nano nginx-deployment.yaml

```
GNU nano 7.2 nginx-deployment.yaml *
apiVersion: apps/v1
kind: Deployment
metadata:
  name: deployment
  labels:
    app: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
          volumeMounts:
            - name: web-content
              mountPath: /usr/share/nginx/html
      volumes:
        - name: web-content
          persistentVolumeClaim:
            claimName: my-pvc
```

Running deployment

kubectl apply -f nginx-deployment.yaml

```
ubuntu@LAPTOP-9M7MBHQ5:~/aws-cli-docker$ kubectl get deployment
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment	2/2	2	2	2m46s

e) Create a [Service](#) associated with the Pod(s) of the HTTP server Deployment.

kubectl expose deployment/deployment --type=LoadBalancer --port=80

```
ubuntu@LAPTOP-9M7MBHQS:~/aws-cli-docker$ kubectl get svc
NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP      PORT(S)
deployment                          LoadBalancer        10.100.175.142   a764d0ebca7904cc89a279402eb1c620-617992065.us-east-1.elb.amazonaws.com  80:32367/TCP
kubernetes                           ClusterIP            10.100.0.1       <none>            443/TCP
my-release-nfs-server-provisioner    ClusterIP            10.100.55.31     <none>            2049/TCP,2049/UDP,32803/TCP,32803/UDP,20048/TCP,20048/UDP,875/TCP,875/UDP,111/TCP,111/UDP,662/TCP,662/UDP 5m33s
```

Details about svc

```
ubuntu@LAPTOP-9M7MBHQS:~/aws-cli-docker$ kubectl describe svc deployment
Name: deployment
Namespace: default
Labels: app=nginx
Annotations: <none>
Selector: app=nginx
Type: LoadBalancer
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.100.175.142
IPs: 10.100.175.142
LoadBalancer Ingress: a764d0ebca7904cc89a279402eb1c620-617992065.us-east-1.elb.amazonaws.com
Port: <unset> 80/TCP
TargetPort: 80/TCP
NodePort: <unset> 32367/TCP
Endpoints: 172.31.14.252:80,172.31.65.28:80
Session Affinity: None
External Traffic Policy: Cluster
Events:
  Type    Reason              Age   From                  Message
  ----    -
  Normal  EnsuringLoadBalancer 3m43s service-controller    Ensuring load balancer
  Normal  EnsuredLoadBalancer 3m39s service-controller    Ensured load balancer
```

Details about nodes

```
ubuntu@LAPTOP-9M7MBHQS:~/aws-cli-docker$ kubectl get nodes -o wide
NAME                                STATUS    ROLES    AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE             KERNEL-VERSION   CONTAINER-RUNTIME
ip-172-31-1-109.ec2.internal        Ready    <none>   27h   v1.31.0-eks-a737599   172.31.1.109   3.215.180.22   Amazon Linux 2     5.10.226-214.880.amzn2.x86_64   containerd://1.7.22
ip-172-31-65-231.ec2.internal       Ready    <none>   27h   v1.31.0-eks-a737599   172.31.65.231  3.83.85.212   Amazon Linux 2     5.10.226-214.880.amzn2.x86_64   containerd://1.7.22
```

Security group rules

Inbound rules <a href="#">Info</a>						
Security group rule ID	Type <a href="#">Info</a>	Protocol <a href="#">Info</a>	Port range <a href="#">Info</a>	Source <a href="#">Info</a>	Description - optional <a href="#">Info</a>	
sgr-0bcba7e5331be4860	Custom TCP	TCP	32367	Custom	Q	<a href="#">Delete</a>
sgr-0d18aca606b337937	HTTP	TCP	80	Custom	Q	<a href="#">Delete</a>
sgr-0a87ee7ac513f066a	NFS	TCP	2049	Custom	Q	<a href="#">Delete</a>
sgr-02fb80fc7780465e6	All traffic	All	All	Custom	Q	<a href="#">Delete</a>
sgr-019ad81915df0308f	All traffic	All	All	Custom	Q	<a href="#">Delete</a>

- f) Create a [Job](#) which mounts the PVC and copies a sample content through the shared NFS PV.

#### Job configuration

```
GNU nano 7.2 job.yaml *
apiVersion: batch/v1
kind: Job
metadata:
  name: nfs-copy-job
spec:
  template:
    spec:
      containers:
      - name: nfs-copy
        image: busybox
        command: ["sh", "-c", "echo '<html><body><h1>Hello, World!</h1></body></html>' > /mnt/index.html"]
        volumeMounts:
        - mountPath: /mnt
          name: web-content
      restartPolicy: Never
      volumes:
      - name: web-content
        persistentVolumeClaim:
          claimName: my-pvc
      backoffLimit: 4
```

#### Running job

```
ubuntu@LAPTOP-9M7MBHQS:~/aws-cli-docker$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
deployment-f8bdfdfc7-qd4wf          1/1     Running   0           27m
deployment-f8bdfdfc7-t24j2          1/1     Running   0           27m
my-release-nfs-server-provisioner-0 1/1     Running   0           30m
nfs-copy-job-q1l5z                  0/1     Completed 0           3m20s
```

- g) Test the HTTP server by showing the sample web content in a browser.

#### Test in console

```
ubuntu@LAPTOP-9M7MBHQS:~/aws-cli-docker$ curl http://3.215.180.22:32367
<html><body><h1>Hello, World!</h1></body></html>
ubuntu@LAPTOP-9M7MBHQS:~/aws-cli-docker$ curl http://3.83.85.212:32367
<html><body><h1>Hello, World!</h1></body></html>
```

#### Test HTTP

##### First replica





## Second replica



Niezabezpieczona | 3.83.85.212:32367

# Hello, World!

Repository: <https://github.com/honeyAsya/LSC>

Command list:

nano ~/.aws/credentials or aws configure

nano Dockerfile

docker build -t aws-cli-docker .

docker run --rm -it aws-cli-docker --version

docker run --rm -it -v ~/.aws:/root/.aws aws-cli-docker -

aws sts get-caller-identity

aws eks update-kubeconfig --region us-east-1 --name Lab5

curl https://raw.githubusercontent.com/helm/helm/main/scripts/get-helm-3 | bash

helm repo add nfs-ganesha-server-and-external-provisioner <https://kubernetes-sigs.github.io/nfs-ganesha-server-and-external-provisioner/>

helm install my-release nfs-ganesha-server-and-external-provisioner/nfs-server-provisioner

nano nfs-pvc.yaml

kubectl apply -f nfs-pvc.yaml

kubectl get pvc

nano nginx-deployment.yaml

kubectl apply -f nginx-deployment.yaml

```
kubectl expose deployment/deployment --type=LoadBalancer --port=80
```

```
kubectl get nodes
```

```
nano job.yaml
```

```
kubectl apply -f job.yaml
```

```
curl http://3.215.180.22:32367
```